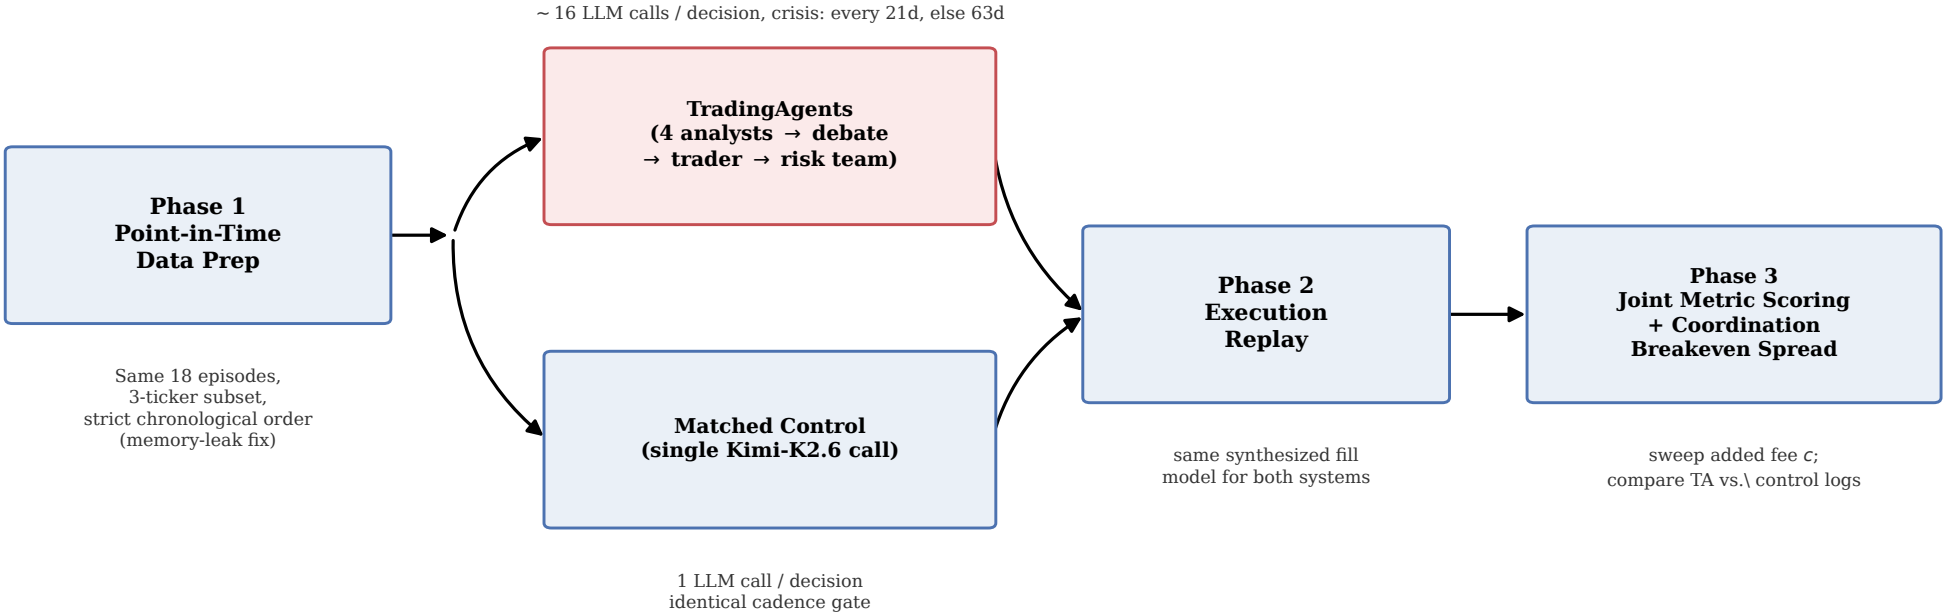
TradingAgents vs. Matched Single-Agent Control Pipeline
~ 16 LLM calls / decision, crisis: every 21d, else 63d
Phase 1
Point-in-Time
Data Prep
Same 18 episodes,
3-ticker subset,
strict chronological order
(memory-leak fix)
TradingAgents
(4 analysts → debate
→ trader → risk team)
Matched Control
(single Kimi-K2.6 call)
1 LLM call / decision
identical cadence gate
Phase 2
Execution
Replay
same synthesized fill
model for both systems
Phase 3
Joint Metric Scoring
+ Coordination
Breakeven Spread
sweep added fee c;
compare TA vs.\ control logs